JINTAO ZHANG

✉ zhang-jt24@mails.tsinghua.edu.cn | ☎ +86 185-1911-5711 (WeChat ID) | 🌐 https://jt-zhang.github.io

EDUCATION

Tsinghua University — **Aug. 2024 – Dec.2026 (Expected)**
Ph.D. of Science in Computer Science and Technology — **Beijing, China**
Supervisor: Prof. Jun Zhu 🔗 Homepage, and Prof. Jianfei Chen 🔗 Homepage.

Tsinghua University — **Aug. 2021 – Jun. 2024**
Master of Science in Computer Science and Technology — **Beijing, China**
Supervisor: Prof. Guoliang Li 🔗 Homepage.

Xidian University — **Sep. 2017 – Jun. 2021**
Bachelor of Science in Computer Science and Technology — **Shaanxi, China**

SELECTED PUBLICATIONS

[PrePrint]
Technical Report
TurboDiffusion: Accelerating Video Diffusion Models by 100-200 Times. [Code (3.1K Stars)]
Jintao Zhang, Kaiwen Zheng, Kai Jiang, Haoxu Wang, Ion Stoica, Joseph E. Gonzalez, Jianfei Chen, Jun Zhu.

[ICLR 2026]
Full Paper, CCF-A
SLA: Beyond Sparsity in Diffusion Transformers via Fine-Tunable Sparse-Linear Attention.
Jintao Zhang, Haoxu Wang, Kai Jiang, ⋯ , Ion Stoica, Joseph E. Gonzalez, Jianfei Chen, Jun Zhu. [Code (0.3K Stars)]

[NeurIPS 2025]
Spotlight, CCF-A
SageAttention3: Microscaling FP4 Attention for Inference and An Exploration of 8-Bit Training.
Jintao Zhang, Jia Wei, Pengle Zhang, Xiaoming Xu, Haofeng Huang, Haoxu Wang, Kai Jiang, Jianfei Chen, Jun Zhu.

[ICML 2025]
Full Paper, CCF-A
SpargeAttention: Accurate Sparse Attention Accelerating Any Model Inference. [Code (0.9K Stars)]
Jintao Zhang, Chendong Xiang, Haofeng Huang, Haocheng Xi, Jia Wei, Jun Zhu, Jianfei Chen.

[ICML 2025]
Full Paper, CCF-A
SageAttention2: Efficient Attention with Thorough Outlier Smoothing and Per-thread INT4 Quantization
Jintao Zhang, Haofeng Huang, Pengle Zhang, Jia Wei, Jun Zhu, Jianfei Chen.

[ICLR 2025]
Full Paper, CCF-A
SageAttention: Accurate 8-Bit Attention for Plug-and-play Inference Acceleration.[Code (3K Stars)]
Jintao Zhang, Jia Wei, Pengle Zhang, Jun Zhu, Jianfei Chen.

[ICDE 2025]
Full Paper, CCF-A
SAGE: A Framework of Precise Retrieval for RAG.
Jintao Zhang, Guoliang Li, Jinyang Su.

[SIGMOD 2024]
Full Paper, CCF-A
PACE: Poisoning Attacks on Learned Cardinality Estimation.
Jintao Zhang, Guoliang Li, Chao Zhang, Chengliang Chai.

[ICDE 2023]
Full Paper, CCF-A
AutoCE: An Accurate and Efficient Model Advisor for Learned Cardinality Estimation.
Jintao Zhang, Chao Zhang, Guoliang Li, Chengliang Chai.

[TKDE 2025]
Full Paper, CCF-A
A Lightweight Learned Cardinality Estimation Model.
Yaoyu Zhu, **Jintao Zhang#** (*corresponding author*), Guoliang Li#, Jianhua Feng.

[VLDB 2022]
Full Paper, CCF-A
Learned Cardinality Estimation: Design Space Exploration and Comparative Evaluation.
Ji Sun*, **Jintao Zhang*** (*equal first contribution*), Zhaoyan Sun, Guoliang Li, Nan Tang.

[PrePrint]
Full Paper
A Survey of Efficient Attention Methods: Hardware-efficient, Sparse, Compact, and Linear Attention
Jintao Zhang, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pengle Zhang, et al.

RESEARCH INTERESTS

- My research focuses on **(1)** Efficient Machine Learning System, specializing in accelerating model training and inference; (2) Data Management, specializing in query optimization and high-quality data acquisition.

SELECTED AWARDS

Tencent Project Up Scholarship (15 globally/year)	2026
Bytedance ScholarShip (20 globally/year)	2025
Siebel Scholars Scholarship (80 globally/year)	2023
China National Scholarship	2019 & 2020 & 2025