JINTAO ZHANG

ZiQiang Technology Building, Tsinghua University, Haidian District, Beijing, China 100084

zhang-jt24@mails.tsinghua.edu.cn | +86 185-1911-5711 (WeChat ID) | https://jt-zhang.github.io

EDUCATION

Tsinghua University — **Aug. 2024 – Dec.2026 (Expected)**
Ph.D. of Science in Computer Science and Technology — **Beijing, China**
Supervisor: Prof. Jun Zhu Homepage, and Prof. Jianfei Chen Homepage.

Tsinghua University — **Aug. 2021 – Jun. 2024**
Master of Science in Computer Science and Technology (GPA: 3.92) — **Beijing, China**
Supervisor: Prof. Guoliang Li Homepage.

Xidian University — **Sep. 2017 – Jun. 2021**
Bachelor of Science in Computer Science and Technology (GPA: 3.85) — **Shannxi, China**

SELECTED PUBLICATIONS

[NeurIPS 2025] SageAttention3: Microscaling FP4 Attention for Inference and An Exploration of 8-Bit Training.
Full Paper, CCF-A — **Jintao Zhang**, Jia Wei, Pengle Zhang, Xiaoming Xu, Haofeng Huang, Haoxu Wang, Kai Jiang, Jun Zhu, Jianfei Chen

[ICML 2025] SpargeAttention: Accurate Sparse Attention Accelerating Any Model Inference. [code (0.7K Stars)]
Full Paper, CCF-A — **Jintao Zhang**, Chendong Xiang, Haofeng Huang, Haocheng Xi, Jia Wei, Jun Zhu, Jianfei Chen.

[ICML 2025] SageAttention2: Efficient Attention with Thorough Outlier Smoothing and Per-thread INT4 Quantization
Full Paper, CCF-A — **Jintao Zhang**, Haofeng Huang, Pengle Zhang, Jia Wei, Jun Zhu, Jianfei Chen.

[ICLR 2025] SageAttention: Accurate 8-Bit Attention for Plug-and-play Inference Acceleration.[code (2K+ Stars)]
Full Paper, CCF-A — **Jintao Zhang**, Jia Wei, Pengle Zhang, Jun Zhu, Jianfei Chen.

[ICDE 2025] SAGE: A Framework of Precise Retrieval for RAG.
Full Paper, CCF-A — **Jintao Zhang**, Guoliang Li, Jinyang Su.

[SIGMOD 2024] PACE: Poisoning Attacks on Learned Cardinality Estimation.
Full Paper, CCF-A — **Jintao Zhang**, Guoliang Li, Chao Zhang, Chengliang Chai.

[ICDE 2023] AutoCE: An Accurate and Efficient Model Advisor for Learned Cardinality Estimation.
Full Paper, CCF-A — **Jintao Zhang**, Chao Zhang, Guoliang Li, Chengliang Chai.

[TKDE 2025] A Lightweight Learned Cardinality Estimation Model.
Full Paper, CCF-A — Yaoyu Zhu, **Jintao Zhang#** (*corresponding author*), Guoliang Li#, Jianhua Feng.

[VLDB 2022] Learned Cardinality Estimation: Design Space Exploration and Comparative Evaluation.
Full Paper, CCF-A — Ji Sun*, **Jintao Zhang*** (*equal first contribution*), Zhaoyan Sun, Guoliang Li, Nan Tang.

[ICMLW 2025] SageAttention2++: A More Efficient Implementation of SageAttention2.
Short Paper — **Jintao Zhang**, Xiaoming Xu, Jia Wei, Haofeng Huang, Pengle Zhang, Chendong Xiang, Jun Zhu, Jianfei Chen.

[PrePrint] A Survey of Efficient Attention Methods: Hardware-efficient, Sparse, Compact, and Linear Attention
Full Paper — **Jintao Zhang**, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pengle Zhang, et al.

RESEARCH STATEMENT

- My research focuses on (1) Efficient Machine Learning System, specializing in accelerating model training and inference; (2) Data Management, specializing in query optimization and high-quality data acquisition.

SELECTED AWARDS

Siebel Scholars Scholarship (35,000 USD, Less than 100 globally/year) — May. 2023
China National Scholarship — 2019 & 2020

PROFESSIONAL EXPERIENCES AND OTHERS

- **Teacher Assistant.** "Programming Fundamentals" in Tsinghua University for three years.
- **Oral Reports.** ICLR 2025, ICDE 2025, SIGMOD 2024, ICDE 2023, and VLDB 2022. TOEFL Score: 99.