

JINTAO ZHANG

✉ zhang-jt24@mails.tsinghua.edu.cn | 📞 +86 185-1911-5711 (WeChat ID) | 🌐 <https://jt-zhang.github.io>

EDUCATION

Tsinghua University Ph.D. in Computer Science and Technology Supervisor: Prof. Jun Zhu 🔗 Homepage , Prof. Jianfei Chen 🔗 Homepage .	Aug. 2024 – Dec. 2026 Beijing, China
Tsinghua University Master in Computer Science and Technology Supervisor: Prof. Guoliang Li 🔗 Homepage .	Aug. 2021 – Jun. 2024 Beijing, China
Xidian University Bachelor of Science in Computer Science and Technology	Sep. 2017 – Jun. 2021 Shaanxi, China

SELECTED PUBLICATIONS

[PrePrint]	TurboDiffusion : Accelerating Video Diffusion Models by 100-200 Times [Code (3.6K Stars)]
Technical Report	Jintao Zhang, Kaiwen Zheng, Kai Jiang, Haoxu Wang, Ion Stoica, Joseph E. Gonzalez, Jianfei Chen, Jun Zhu.
[ICLR 2026]	SLA : Beyond Sparsity in Diffusion Transformers via Fine-Tunable Sparse-Linear Attention
Full Paper, CCF-A	Jintao Zhang, Haoxu Wang, Kai Jiang, ·, Ion Stoica, Joseph E. Gonzalez, Jianfei Chen, Jun Zhu. [Code (0.3K Stars)]
[NeurIPS 2025]	SageAttention3 : Microscaling FP4 Attention for Inference and An Exploration of 8-Bit Training
Spotlight, CCF-A	Jintao Zhang, Jia Wei, Pengle Zhang, Xiaoming Xu, Haofeng Huang, Haoxu Wang, Kai Jiang, Jianfei Chen, Jun Zhu.
[ICML 2025]	SpargeAttention : Accurate Sparse Attention Accelerating Any Model Inference [Code (1K Stars)]
Full Paper, CCF-A	Jintao Zhang, Chendong Xiang, Haofeng Huang, Haocheng Xi, Jia Wei, Jun Zhu, Jianfei Chen.
[ICML 2025]	SageAttention2 : Efficient Attention with Thorough Outlier Smoothing and Per-thread INT4 Quantization
Full Paper, CCF-A	Jintao Zhang, Haofeng Huang, Pengle Zhang, Jia Wei, Jun Zhu, Jianfei Chen.
[ICLR 2025]	SageAttention : Accurate 8-Bit Attention for Plug-and-play Inference Acceleration [Code (3.5K Stars)]
Full Paper, CCF-A	Jintao Zhang, Jia Wei, Pengle Zhang, Jun Zhu, Jianfei Chen.
[ICDE 2025]	SAGE : A Framework of Precise Retrieval for RAG
Full Paper, CCF-A	Jintao Zhang, Guoliang Li, Jinyang Su.
[SIGMOD 2024]	PACE : Poisoning Attacks on Learned Cardinality Estimation
Full Paper, CCF-A	Jintao Zhang, Guoliang Li, Chao Zhang, Chengliang Chai.
[ICDE 2023]	AutoCE : An Accurate and Efficient Model Advisor for Learned Cardinality Estimation
Full Paper, CCF-A	Jintao Zhang, Chao Zhang, Guoliang Li, Chengliang Chai.
[TKDE 2025]	A Lightweight Learned Cardinality Estimation Model
Full Paper, CCF-A	Yaoyu Zhu, Jintao Zhang [#] (<i>corresponding author</i>), Guoliang Li [#] , Jianhua Feng.
[VLDB 2022]	Learned Cardinality Estimation : Design Space Exploration and Comparative Evaluation
Full Paper, CCF-A	Ji Sun*, Jintao Zhang* (<i>equal first contribution</i>), Zhaoyan Sun, Guoliang Li, Nan Tang.
[PrePrint]	A Survey of Efficient Attention Methods : Hardware-efficient, Sparse, Compact, and Linear Attention
Full Paper	Jintao Zhang, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pengle Zhang, et al.
[PrePrint]	Speculative Decoding for Autoregressive Video Generation
Full Paper	Yuezhou Hu, Jintao Zhang ^{*#} (<i>equal first contribution and corresponding author</i>).

RESEARCH INTERESTS

- My research focuses on (1) Efficient Machine Learning System, specializing in accelerating model training and inference; (2) Real-time Video Generation, covering the full pipeline from data preparation to model training.

SELECTED AWARDS

Tencent Project Up Scholarship (15 globally/year)	2026
Bytedance ScholarShip (20 globally/year)	2025
Siebel Scholars Scholarship (80 globally/year)	2023
Zhang Bo Artificial Intelligence Scholarship (2/year)	2026